

# A Reconfigurable Streaming Deep Convolutional Neural Network Accelerator for Internet of Things

Li Du, *Member, IEEE*, Yuan Du, *Member, IEEE*, Yilei Li, Junjie Su, Yen-Cheng Kuan, *Member, IEEE*, Chun-Chen Liu, and Mau-Chung Frank Chang, *Fellow, IEEE*

**Abstract**—Convolutional neural network (CNN) offers significant accuracy in image detection. To implement image detection using CNN in the Internet of Things (IoT) devices, a streaming hardware accelerator is proposed. The proposed accelerator optimizes the energy efficiency by avoiding unnecessary data movement. With unique filter decomposition technique, the accelerator can support arbitrary convolution window size. In addition, max-pooling function can be computed in parallel with convolution by using separate pooling unit, thus achieving throughput improvement. A prototype accelerator was implemented in TSMC 65-nm technology with a core size of 5 mm<sup>2</sup>. The accelerator can support major CNNs and achieve 152GOPS peak throughput and 434GOPS/W energy efficiency at 350 mW, making it a promising hardware accelerator for intelligent IoT devices.

**Index Terms**—Convolution neural network, deep learning, hardware accelerator, IoT.

## I. INTRODUCTION

**M**ACHINE Learning offers many innovative applications in the IoT devices, such as face recognition, smart security and object detection [1]–[3]. State-of-the-art machine-learning computation mostly relies on the cloud servers [4], [5]. Benefiting from the graph processing unit (GPU)’s powerful computation ability, the cloud can process high throughput video data coming from the devices and use CNN to achieve unprecedented accuracy on most AI applications [6]. However, this approach has its own drawbacks. Since the network connectivity is necessary for cloud-based AI applications, those applications cannot run in the areas where there is no network coverage. In addition, data transfer through network induces significant latency, which is not acceptable for real-time AI applications such as security

system. Finally, most of the IoT applications have a tough power and cost budget which could tolerate neither local GPU solutions nor transmitting massive amounts of image and audio data to data center servers [7].

To address these challenges, a localized AI processing scheme is proposed. The localized AI processing scheme aims at processing the acquired data at the client side and finishes the whole AI computation without communication network access. Conventionally, this is done through local GPU or DSP. However, this results in a limited computation ability and relatively large power consumption, making it not suitable for running computation-hungry neural network such as CNN on power limited IoT devices [8]. Consequently, it is crucial to design a dedicated CNN accelerator inside the IoT devices that can support a high performance AI computation with minimal power consumption. Some of the reported works in the neural network acceleration are focusing on providing an architecture for computing general neural network. For example, in [9], an efficient hardware architecture is proposed based on the sparsity of the neural network through pruning the network properly. However, it is a more general architecture to compute the fully-connected deep neural network without considering parameter reuse. On the contrary, the CNN has its unique feature that the filters’ weights will be largely reused throughout each image during scanning. Benefiting from this feature, many dedicated CNN hardware accelerators are reported [10]–[12]. Most of reported CNN accelerators only focus on accelerating the convolution part while ignoring the implementation of the pooling function, which is a common layer in the CNN network. In [10], a CNN hardware accelerator using a spatial architecture with 168 processing elements is demonstrated. In [11], another dedicated convolution accelerator with loop-tiling optimization is reported. Since pooling function is not implemented in those accelerators, the convolution results must be transferred to CPU/GPU to run pooling function and then fed back to the accelerator to compute the next layer. This data movement not only consumes much power but also limits overall performance. On the other hand, some works report highly configurable neural network processors but they require complicated data flow control. This adds hardware overhead to IoT devices. For example, [12] reports a CNN processor occupying 16 mm<sup>2</sup> silicon area in 65nm CMOS technology, which can be intolerable for low-cost IoT chips. In addition, several recent reports, such as [13], proposed to use memristors to perform neuromorphic computing for CNN. However, the fabrication of memristors currently is still not supportive in major

Manuscript received December 8, 2016; revised June 4, 2017; accepted July 27, 2017. Date of publication August 16, 2017; date of current version January 5, 2018. This paper was recommended by Associate Editor A. Sangiovanni Vincentelli. (Corresponding authors: Li Du; Yen-Cheng Kuan; Mau-Chung Frank Chang.)

L. Du, Y. Du, and C.-C. Liu are with the High Speed Electronics Laboratory, University of California at Los Angeles, Los Angeles, CA 90095 USA, and also with Kneron Inc., San Diego, CA 92121 USA (e-mail: dl1989113@ucla.edu).

Y. Li is with the Novumind Inc., Santa Clara, CA 95054 USA.

J. Su is with Kneron Inc., San Diego, CA 92121 USA.

Y.-C. Kuan is with National Chiao Tung University, Hsinchu 30010, Taiwan (e-mail: yckuan@g2.nctu.edu.tw).

M.-C. F. Chang is with the High Speed Electronics Laboratory, University of California at Los Angeles, Los Angeles, CA 90095 USA, and also with National Chiao Tung University, Hsinchu 30010, Taiwan (e-mail: mfchang@ee.ucla.edu).

Color versions of one or more of the figures in this paper are available online at <http://ieeexplore.ieee.org>.

Digital Object Identifier 10.1109/TCSI.2017.2735490

CMOS foundries [14]. Thus this architecture is hard to embed into the IoT chips.

In this study, we propose a new streaming hardware architecture for CNN inference at the IoT platform and assume the CNN model is pre-trained. We focus on the optimization of the data-movement flow to minimize data access and achieve high energy efficiency for computation. A new methodology is also proposed to decompose large kernel-sized computation to many parallel small kernel-sized computations. Together with the integrated pooling function, our proposed accelerator architecture can support completed one-stop CNN acceleration with both arbitrarily sized convolution and reconfigurable pooling. The main contribution of this paper includes:

- 1) A CNN accelerator design using streaming data flow to achieve optimal energy efficiency.
- 2) An interleaving architecture to enable parallel computing for multiple output features without SRAM input bandwidth increment.
- 3) A methodology to decompose large-sized filter computation into many small-sized filter computation, achieving high reconfigurability without adding additional hardware penalty.
- 4) A supplementary pooling block that can support pooling function while the main engine serves for CNN computation.
- 5) A prototype design with FPGA verification, which can achieve a peak performance of 152 GOPS and energy efficiency of 434 GOPS/W.

The paper is organized as follows. In Section II, we first introduce the main layers composing CNNs. In Section III, we introduce our system's overview architecture. In Section IV, we discuss the proposed streaming architecture to achieve high-efficiency convolution computation, filter decomposition technique to provide reconfigurability and pooling implementation. Key modules' design is explained in Section V. Finally, the experimental results are reported in Section VI and the conclusion is drawn in Section VII.

## II. LAYER DESCRIPTION

The state-of-art CNN networks (e.g., AlexNet, VGG-18, etc.) [15]–[17] are mainly composed of three typical layers: convolution layer, pooling layer and classification layer. Convolution layer composes the majority of the neural network, with pooling layer inserted between two convolution layers to achieve intermediate data size reduction and non-linear mapping. Classification layer is usually included as the last layer of the CNN, which does not require a large amount of computation. Here, we assume that the classification layer can be realized through software computation and will not be implemented in the hardware accelerator.

The following subsection will explain the convolution layer and the pooling layer's functions in details.

### A. Convolution Layer

The primary role of a convolution layer is to apply convolution function to map the input (previous) layer's images to the next layer.




Fig. 1. Example of computation of a CNN layer.

Since each input layer can have multiple input features (referred as channels afterward), the convolution is 3D. Unlike regular convolution, where it took the whole input data to generate one output data, the convolution in a neural network is localized through forming a regional filter window in each individual input channel. This set of the regional filter windows is regarded as one filter. The output data is obtained through computing the inner product of the filter weight and the input data covered by the filter. An output feature can be obtained by using the convolution filter to scan the input channels. Multiple output features can be computed by using different filters. In addition, a separated bias weight will be added in each final filtered result. The arithmetical representation of this function is shown as (1).

$$\begin{aligned}
 O[io][r][c] &= B[io] + \sum_{ii=0}^{Fi-1} \sum_{i=0}^{K-1} \sum_{j=0}^{K-1} \\
 &\quad \times I[ii][s \times r + i][s \times c + j] \times W[io][ii][i][j] \\
 0 \leq io < Fo, \quad 0 \leq ii < Fi, \quad 0 \leq r < ROW, \quad 0 \leq c < COL
 \end{aligned} \tag{1}$$

Here,  $io$  represents the current output-feature's index number,  $Fi$  and  $Fo$  represents the total number of the input channels and output features.  $r$  and  $c$  represents the current output-feature's data's row and column number;  $s$  is the stride size of the convolution window,  $W$  represents the filter weight and  $B$  represents the bias weight of each filter.  $K$ ,  $ROW$  and  $COL$ , are the kernel size, output-feature row size and column size respectively.

With the above parameters' definition, the input layer has  $Fi$  channels. Each channel's width is  $COL \times s$  and height is  $ROW \times s$ . The layer output includes  $Fo$  features. Each feature's width is  $COL$  and height is  $ROW$ . Filter number is same as the output-feature number. In each filter, it is constructed through  $Fi$  separated filter window. Each window's kernel size is  $K$ . The overall convolution procedure is represented as Fig. 1.

### B. Pooling Layer

In addition to the convolution layer, pooling layer is also an important part of the regular CNN. The role of the pooling layer is to extract information from a set of neighboring image pixels in each channel. Typically, the pooling layer can be separated into two categories: max pooling layer and average pooling layer. The max pooling layer selects the maximum image data's value within the pooling window, while the average pooling layer provides the average value of the data within the pooling window. The mathematical representations




Fig. 2. Example of computation of a max pooling layer.

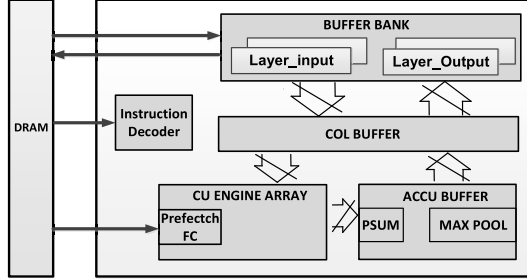


Fig. 3. Overall architecture of the CNN accelerator.

of these two pooling operations are defined as (2) and (3). Fig. 2 is an example of the max pooling function.

$$O_{avg}[r][c] = avg \begin{bmatrix} I[r][c] & \cdots & I[r][c+K-1] \\ \vdots & \ddots & \vdots \\ I[r+K-1][c] & \cdots & I[r+K-1][c+K-1] \end{bmatrix} \quad (2)$$

$$O_{max}[r][c] = max \begin{bmatrix} I[r][c] & \cdots & I[r][c+K-1] \\ \vdots & \ddots & \vdots \\ I[r+K-1][c] & \cdots & I[r+K-1][c+K-1] \end{bmatrix} \quad (3)$$

Here  $I[r][c]$  represents the input channel's data at the position (r,c) and the kernel size of the pooling window is K.

### III. SYSTEM OVERVIEW

The overall streaming architecture of the CNN accelerator is shown in Fig. 3. It is already proved that deep networks can be represented with 16-bit fixed-point number with stochastic rounding and incur little to no degradation in the classification accuracy [18]. In addition, an implementation of the 16-bit floating adder costs much more logic gates compared to that of the 16-bit fixed-point adder [8]. Thus, the data format of this accelerator is set as the 16-bit fixed point. The floating weights will be truncated to the 16-bit fixed point during the training. The accelerator includes a 96 Kbyte single port SRAM as the buffer bank to store the intermediate data and exchange data with the DRAM. The buffer bank is separated into two sets. One for the input data of the current layer and the other one is to store the output data. The input channels and output features are numbered. In each set, the buffer bank is further divided into Bank A and Bank B. Bank A is used to store the odd-number channels/features, while Bank B is used to store the even-number channels/features. In addition, a COL BUFFER module is implemented to remap the buffer

bank's output to the convolution unit (CU) engine's input. The CU engine is composed of sixteen convolution units to enable highly parallel convolution computation. Each unit can support the convolution with a kernel size up to three. A pre-fetch controller is included inside the engine to periodically fetch the parameters from Direct Memory Access (DMA) controller and update the weights and bias values in the engine. Finally, an accumulation (ACCU) buffer with scratchpad is implemented in the accelerator. The scratchpad is used together with the accumulator to accumulate and store the partial convolution results coming from the CU engine. A separated max pooling module is also embedded in the ACCU buffer to pool the output-layer data if necessary.

The control of this accelerator is through 16-bit Advanced Extensible Interface (AXI) bus, the command decoder is integrated inside the accelerator. The commands for the processed CNN net are pre-stored in the DRAM in advance, and will be automatically loaded to a 128-depth command FIFO when the accelerator is enabled.

The commands can be divided into two categories: configuration commands and execution commands. Configuration commands are inserted between multiple layers to configure the upcoming layer's property, such as channel size and numbers, enable ReLU function or max pooling function. The execution commands are to initiate the convolution/pooling computation. The configuration of the shifting address value for large-sized convolution filter is also included in the execution commands (explained in Section V).

The convolution begins with resetting the image scratchpad in the ACCU buffer. Then the input-layer data will be sent to CU engine sequentially. The CU engine will calculate the inner product of each channel's data with its corresponding output feature's filter's weight. Output results from the CU engine will be passed to the ACCU Buffer block and accumulated with the stored results in the scratchpad. After all the channels are scanned, the accumulated image in the scratchpad will be sent back to the Buffer bank as one of the output features.

After finishing the computation of the 1<sup>st</sup> feature, the CNN accelerator will duplicate the convolution procedure described above with updated filter weights from the DRAM, to generate the next output feature. This procedure will be continuously reproduced till all the features are calculated. The overall diagram showing this procedure is drawn as Fig. 4.

### IV. STREAMING AND RECONFIGURABLE FEATURES

The proposed CNN accelerator achieves reconfigurability and high energy efficiency through using three techniques below:

1. Using filter decomposition technique to support large kernel-sized filter's computation through using only  $3 \times 3$ -sized computation unit.
2. Streaming data flow to minimize bus control and module interface, thus reducing hardware cost while achieving high energy efficiency.
3. Separate pooling blocks to compute max pooling in parallel with convolution and reuse the convolution engine for average pooling functions to achieve minimum hardware design cost.






Fig. 4. Convolution Computation Procedure. Input-layer data is stored by channels in the buffer bank and will be fed to the CU engine sequentially. The weight is stored in DRAM and will be fed to the CU engine during the convolution.




Fig. 5. An  $5 \times 5$  Filter decomposed into four  $3 \times 3$  sub filter. F0, F1, F2, F3's shift address are (0,0), (0,3), (3,0), (3,3).

#### A. Filter Decomposition

The filter's kernel size in a typical CNN network can range from very small size ( $1 \times 1$ ) to very large size ( $11 \times 11$ ). Hardware convolution engine is usually designed for a certain kernel size and can only support filter computation below its limited size. So when computing the convolution with kernel size above its limitation, the accelerator needs to either leave the software to do the computation or add additional hardware unit for large kernel-sized filter convolution.

To minimize the hardware resource usage, a filter decomposition algorithm is proposed to compute any large kernel-sized ( $> 3 \times 3$ ) convolution through using only  $3 \times 3$ -sized CU. The algorithm begins with examining the kernel size of the filter. If the original filter's kernel size is not an exact multiple of three, zero padding weights will be added in the original filter's kernel boundary to extend the original filter's kernel size to be a multiple of three. Because the added weights in the boundary are 0, so the extended filter will result in same output value compared with the original filter during the computation. Next, the extended filters will be decomposed into several  $3 \times 3$ -sized filters. Each filter will be assigned a shift address based on its top left weight's relative position in the original filter. For example, Fig. 5 is an example of decomposing a  $5 \times 5$  filter into four  $3 \times 3$  filters. One row and column zero padding are added in the original filter. The decomposed filters: F0, F1, F2, F3's shift address are (0,0), (0,3), (3,0), (3,3).

After that, we compute each decomposed filter with the input layer separately, generating several decomposed output features. Finally, we recombine those decomposed features into one final output feature through (4).

$$I_o(X, Y) = \sum_i I_{d_i}(X + x_i, Y + y_i) \quad (4)$$

Here,  $I_o$  represents the output image,  $I_{d_i}$  represents  $i$ 's decomposed filter's output image,  $(X, Y)$  represents the current output data's coordinate address,  $(x_i, y_i)$  represents  $i$ 's filter's shift address.

The arithmetical derivation of this filter decomposition can be described as (5)

$$\begin{aligned} F_{3K}(a, b) &= \sum_{i=0}^{3K-1} \sum_{j=0}^{3K-1} f(i, j) \times I_i(a + i, b + j) \\ &= \sum_{i=0}^{K-1} \sum_{j=0}^{K-1} \sum_{l=0}^2 \sum_{m=0}^2 f(3i + l, 3j + m) \\ &\quad \times I_i(a + 3i + l, b + 3j + m) \\ &= \sum_{i=0}^{K-1} \sum_{j=0}^{K-1} F_{3-i-j}(a + 3i, b + 3j) \end{aligned} \quad (5)$$

$$\begin{aligned} F_{3-i-j}(a, b) &= \sum_{m=0}^2 \sum_{l=0}^2 f(3i + l, 3j + m) \\ &\quad \times I_i(a + 3i + l, b + 3j + m) \end{aligned}$$

$$0 \leq i < K - 1; \quad 0 \leq k < K - 1; \quad (6)$$

Here  $F_{3K}(a, b)$  represents a filter with kernel size of  $3K$  and its top-left weight is multiplied with the pixel's value at the position  $(a, b)$  in the image. Each weight in the filter is represented as  $f(i, j)$  where the  $(i, j)$  represents the weight's position relative to the top-left weight inside the filter and  $I_i(a + 3i + l, b + 3j + m)$  represent the image pixel's value at the position of  $(a + 3i + l, b + 3j + m)$  in the image.  $F_{3-i-j}$  represents  $K^2$ 's different  $3 \times 3$  kernel-sized filter with its computation function defined as (6). In addition, the  $3i$  and  $3j$  can represent as the shifting address of each  $3 \times 3$  filter.

Based on (5) and (6), we can approve that a  $3K \times 3K$  filter's computation can be decomposed into  $K^2$  different  $3 \times 3$  filters' calculation without any loss of the computation accuracy. Fig. 6 is an example of using this filter decomposed technique to compute a  $5 \times 5$  convolution.

This decomposition technique provides a benefit of maximized hardware resource usage at the penalty of adding additional zero padding in the filter boundary. Although this added zero padding results in a waste of the computation resource, the overall efficiency loss is relatively small in the CNN net. On the contrary, the CU engine design becomes much simpler as it only needs to support convolution filter size of  $1 \times 1$  and  $3 \times 3$ . The overall efficiency loss can be computed based on the (7).

$$EL = \frac{MAC_{zero\_padding}}{MAC_{total}} \quad (7)$$

Here the  $MAC_{total}$  represented the total multiply-accumulate operation(MAC) the engine takes to compute a CNN network, while  $MAC_{zero\_padding}$  represented the MAC operation used to compute the zero-padding part. For example, a  $11 \times 11$  filter actually has  $\frac{23}{144}$  MAC operation used on computing zero-padding part, resulting in an efficiency loss of 16%.

Table I is a comparison of different major CNN nets efficiency loss by using this decomposition technique.

As TABLE I shows, AlexNet exhibits the largest efficiency loss since it has a large  $11 \times 11$  filter in the first layer. On the contrary, small filter-sized nets such as Resnet-18, Resnet-50, Inception V3, have a very small efficiency loss due to zero-padding. Hence, they are well-suited for this architecture.

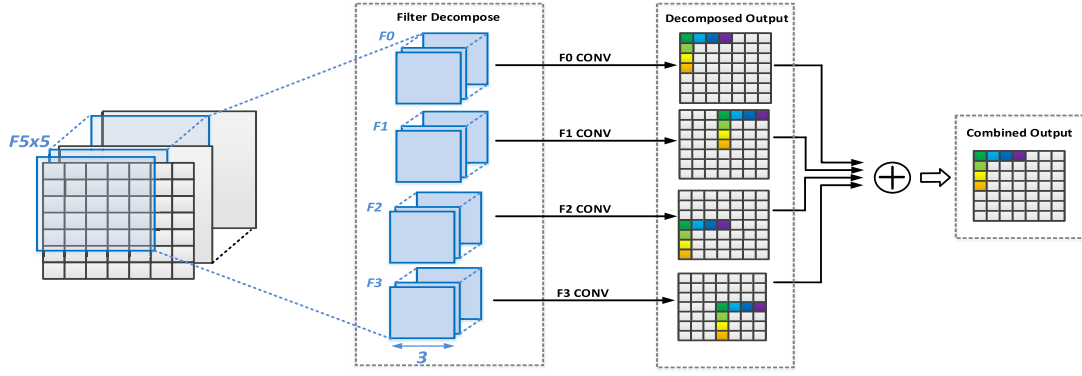


Fig. 6. Filter decomposition technique to compute a  $5 \times 5$  filter on the  $7 \times 7$  image. The Filter is decomposed into F0, F1, F2, F3, generating four sub-images. The sub-images are summed based on their filter's shift address. Same color's pixels will be added together to generate the corresponding pixels in the output image.

TABLE I  
CONVOLUTION EFFICIENCY LOSS THROUGH  
DECOMPOSITION TECHNIQUE

Net Type	Filter Kernel Size	Efficiency Loss
AlexNet	3-11	13.74%
ResNet-18	1-7	1.64%
ResNet-50	1-7	0.12%
Inception V3	1-5	0.89%

### B. Streaming Architecture

To minimize the data movement and achieve optimal energy efficiency for the convolution computation, a streaming architecture is proposed for the CNN accelerator. For a regular CNN convolution, it includes multiple levels of data and weights reuse:

1. Every set of the filter weights is reused to scan the whole channel's image.
2. Every output feature is generated through scanning the same input layer.

The streaming architectures reduce the data movement through benefiting the above-listed features in CNN convolution.

1) *Filter Weight Reuse*: In each filter, the weights between kernels are different. Each kernel's weights will only be used with the particular input channel's data. To benefit from this, all the filter weights are stored in the DRAM and will only be fetched into the accelerator during the convolution.

During the  $3 \times 3$  convolution, the fetched filter weights will be stored in the CU engine and input channel's image data will stream into the CU engine. The CU engine will produce the inner product between the weights and the streamed-in data, generating a corresponding output feature's partial result to the ACCU buffer for accumulation. The weights in the CU engine will not be updated until the whole channel is scanned. The  $1 \times 1$  convolution follows the similar approach as the  $3 \times 3$  convolution except that seven out of nine multipliers are turned off during the convolution. The left two multipliers will be turned on to calculate two different output features' partial summation result simultaneously.

Fig. 7 is an example showing one filter window movement of this flow. The real implementation includes sixteen

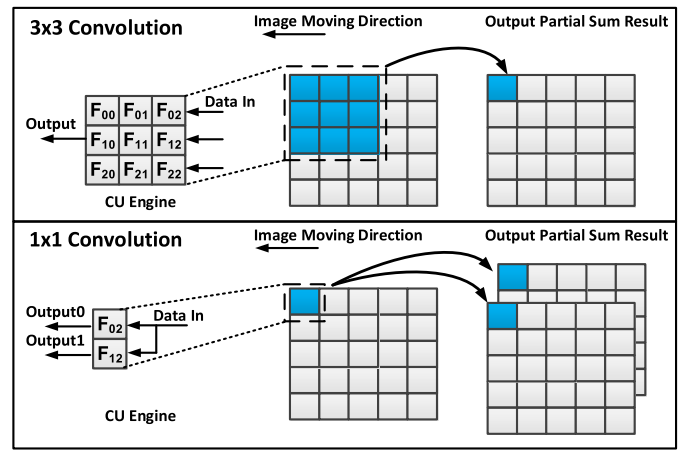


Fig. 7. Data flow of the streaming architecture.

$3 \times 3$  filter windows to process multiple rows' data simultaneously. By using this filter window to scan the input channel, the data flow and module interface become much simpler and hence the hardware design cost is reduced.

The output bandwidth of the Buffer Bank is set to be 256 bits/cycle with each data size as 16 bits, corresponding to stream in sixteen data from different rows to the CU engine simultaneously. The sixteen data are divided into two sets: eight data are from the odd-number channels and the other eight data is from the even-number channels.

To maximize the usage of the buffer bank's output bandwidth, a two-rows' FIFO buffer is paired with each set of the row data, transferring the eight input rows to ten overlapping output rows. This enables running eight  $3 \times 3$  CU in parallel for each set of the row data. The FIFO buffer included in the COL buffer is shown in Fig. 8. Here we only draw half sized COL buffer for the even-number channels' data. Real implementation includes the FIFO buffer for both even and odd channels.

2) *Input Channel Reuse*: In the  $1 \times 1$  convolution, each output-feature data computation only requires one multiplication in each channel. This results in wasting a majority of hardware resource as most of the multipliers in the CU engine will not be used.

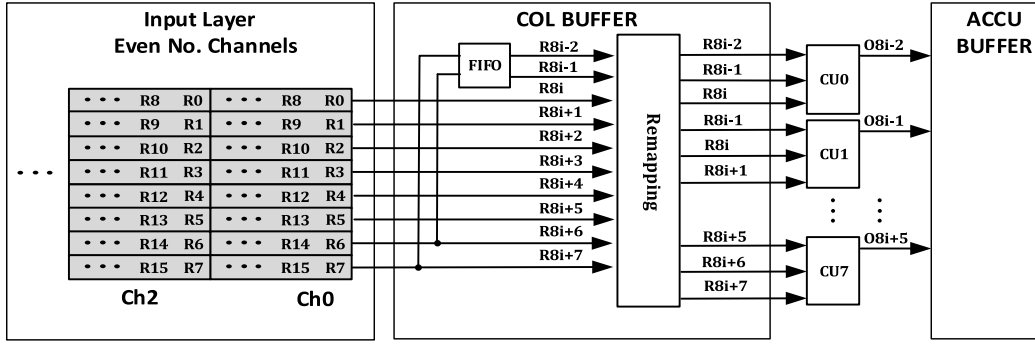


Fig. 8. Half of COL Buffer Architecture of the CNN accelerator, input channels are stored sequentially in the buffer. Each channel is stored by rows,  $R_i$  represents  $i$ 's row's data.

To accelerate the computation in the  $1 \times 1$  convolution, an interleaving architecture is proposed to compute two output features' results in parallel in the  $1 \times 1$  convolution. Since computation of each output feature requires scanning the same input layer, the accelerator can compute multiple output features simultaneously during one scanning. However, if multiple features are computed simultaneously, it will result in a proportional output bandwidth increment for the CU engine. For example, output two features simultaneously will lead the CU engine to generate twice output data bandwidth compared to its input data bandwidth.

To prevent this, an interleaving architecture is proposed. The previous layer's even-number features and odd-number features are stored into two different banks and are used as the input channels in the current layer. Each bank's output bandwidth is 128 bits, corresponding to 8 pixels' data. These two banks' data are fed into the CU simultaneously. Thus, the CU's input data are separated into even-number channel's data and odd-number channel's data. These two sets of data are individually multiplied with two different features' weights, resulting in a total of 32 data (two output features' partial results) at the CU engine's output. However, since the 32 data streams are generated from two different channels, a summation function is required at the CU output to combine the same feature's partial results from different channels. By doing this, the data bandwidth is reduced by half and hence the final output bandwidth of the adder will be same as the input data bandwidth. The detailed implementation of this function is drawn in Fig. 9. Here the  $X(O, 0) \dots X(O, 7)$  represents the 1<sup>st</sup> to the 8<sup>th</sup> row's data of the odd-number channels and the  $X(E, 0) \dots X(E, 7)$  represents the 1<sup>st</sup> to the 8<sup>th</sup> row's data of the even-number channels.  $O(0, 1)$ ,  $E(0, 1)$  are the 1<sup>st</sup> features' partial results from odd-number and even-number channels and  $O(0, 2)$ ,  $E(0, 2)$  are the 2<sup>nd</sup> features' partial results from the odd-number and even-number channels.

### C. Pooling

Pooling functions are also implemented in the accelerator. The pooling functions can be separated into two categories: max pooling and average pooling.

1) *Average Pooling*: To minimize hardware cost, the average pooling function is implemented through reusing the convolution engine. This can be achieved through replacing the

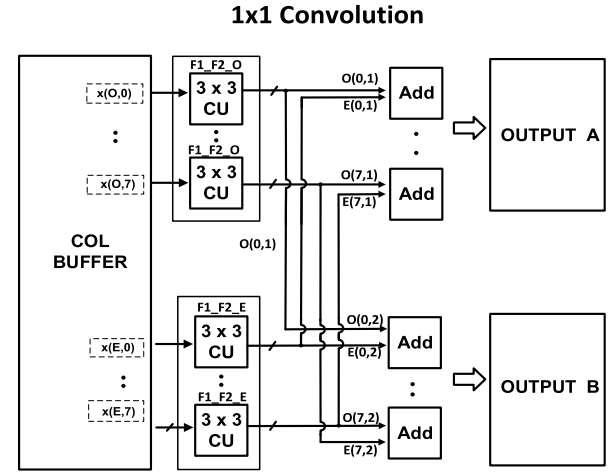


Fig. 9. Data flow of the streaming architecture in the  $1 \times 1$  convolution mode.

average pooling layer with the same kernel-sized convolution layer using the following steps:

1. Create a convolution layer with the output features' number to be equal to the input channels' number. The kernel size is same as the pooling window size
2. In each filter, set the corresponding channel's filter's weight to  $\frac{1}{K^2}$ , where  $K$  is the kernel size. All other channels' filter weights are set to 0.

The arithmetical representation of this convolution layer can be derived as (8).

$$O[io][r][c] = \sum_{ii=0}^I \sum_{i=0}^{K-1} \sum_{j=0}^{K-1} \times I[io][ii][r+i][c+j] \times W[io][ii][i][j]$$

$$W[io][ii][i][j] = \begin{cases} \frac{1}{K^2} & \text{if } ii = io \\ 0 & \text{if } ii \neq io \end{cases} \quad (8)$$

Here  $ii$  and  $io$  are the input-channel number and output-feature number,  $r$  and  $c$  are the output feature's row and column's number,  $W$  represents the weight of the filter.  $K$  is the kernel size of the average pooling window.

2) *Max Pooling*: The max pooling layer is implemented as a separate block inside the ACCU buffer and it is used to pool the output feature coming from the convolution block.

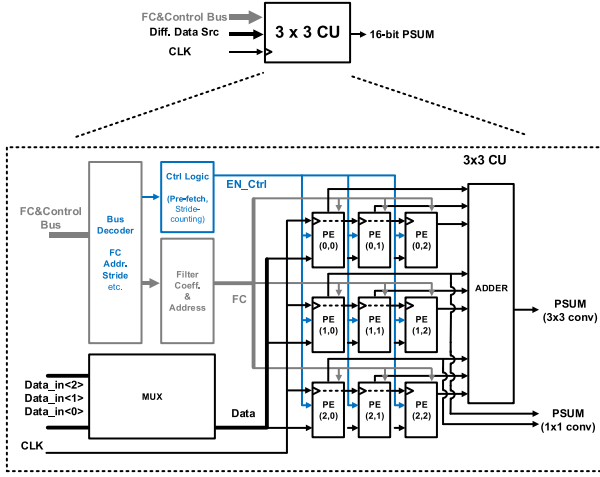


Fig. 10. Implementation of the  $3 \times 3$  CU engine.

The pooling block is designed to support pooling window size of two and three, which covers major CNNs [15]–[17]. The detailed implementation of this block and its connection to the scratchpad will be described in Section V.

## V. MODULE IMPLEMENTATION

In this sections, three major modules: CU engine, ACCU buffer and Max pooling in this accelerator will be discussed.

### A. CU Engine

As described in Section IV, the accelerator uses nine multipliers to form a CU and sixteen CUs to compose a CU engine. The module implementation of the CU is shown in Fig. 10.

The CU engine array includes nine processing engines (PE) and an adder to combine the output. The PE provides a multiplication function for its input data and the filter's weight and meanwhile passes its input data to the next stage's PE through a D flip-flop. The multiplication function can be turned on/off based on the EN\_Ctrl signal to save the computation power when convolution stride size is larger than one.

In the  $3 \times 3$  convolution, the multiplied result will send to the adder in the CU to perform the summation and deliver the summed result to the final output. Filter weights will be fetched from the DRAM through the DMA controller and pre-stored in the CU through a global bus. When one channel is scanned, a synchronized filter updated request signal will be sent to the CU to update the filter weights at the PE's input for the upcoming channel.

In the  $1 \times 1$  convolution, only PE (1,0) and PE (2,0) will be turned on. The adder will be disabled and the two output results are directly fed out as the two output-feature partial results.

### B. ACCU Buffer

The ACCU Buffer is used to accumulate the output partial summation results from the CU engine and meanwhile temporary store the feature output data in its scratchpad, waiting for the buffer bank to read back. The ACCU buffer includes a ping-pong buffer as the scratchpad, an accumulator to sum the partial result, a separate pooling block for max pooling and

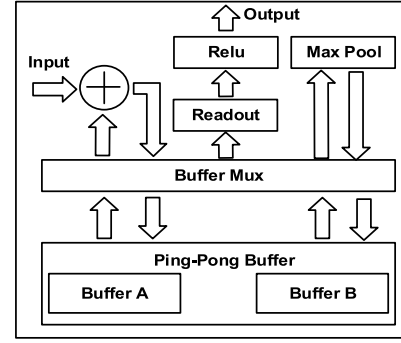


Fig. 11. The ACCU Buffer includes a Ping-Pong buffer formed by the Buffer A and the Buffer B. The two buffers will be switched back and forth between the accumulator and the Readout/Max pool blocks to enable parallel processing.

a readout block to read data from the scratchpad back to the buffer bank.

The ping-pong buffer is separated into two different sub-buffers. During the convolution, only one buffer will be pointed to the accumulator while the other buffer will be connected to the pooling blocks and the readout blocks. This enables the core to process the pooling functions and the convolution functions simultaneously. In addition, reading data from the scratchpad back to the buffer bank can also be processed in parallel with the convolution.

When the accumulator finished accumulating one output feature, the ping-pong buffer will switch its sub-buffers directions, pointing the buffer that stores the output feature to the pooling blocks and the readout blocks. Meanwhile, the sub buffer which previously connected to the pooling side will turn to the accumulator to continuously accumulate the next output-feature partial summation result. In addition, the ReLU function is implemented during the readout. The ReLU function can be realized through zeroing the negative output from the readout blocks.

Compared with the convolution, the readout and pooling functions only need to scan one output feature each time, resulting a much shorter time to process. Benefited from this, the accelerator can continuously run convolution without any speed loss on the pooling and the data readout. The detailed implementation of the ACCU Buffer architecture is shown as Fig.11.

### C. Max Pool

Fig. 12 shows an overview architecture of the max pooling module and its connection to the scratchpad. The scratchpad stored eight rows' data from one output feature in parallel. The eight rows' data share one column address and can be accessed simultaneously. Because of the stride size's difference in the convolution, data stored in the scratchpad may not be all validated. For example, when the stride is equal to 2, only R0, R2, R4, R6 store the validate data. In addition, the pool window's kernel size can also be configured to be 2 or 3.

To accommodate different convolution strides and pool-size cases, a MUX is put in front of the max pooling module to select the validated input data to the corresponding max-pool units. The max-pool unit is implemented with a four-input



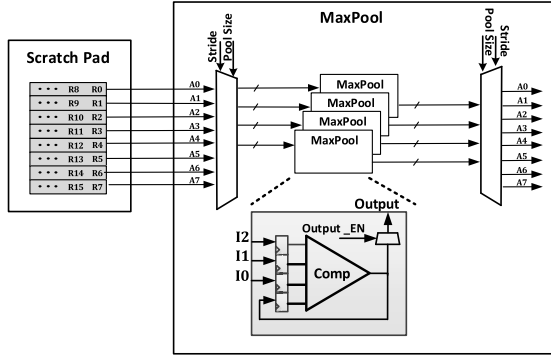


Fig. 12. Overall architecture of the Maxpooling module, Ri represents row i's data. The pooled output will be fed back to the scratchpad.




Fig. 13. Layout view of the accelerator.

comparator and a feedback register to store the intermediate comparator output result. In addition, an internal buffer is embedded in the max pooling module. This is to buffer the intermediate results if some of the data inside the pooling window are not ready.

When a pooling begins, the comparator first takes three input data coming from nearby rows (two data in  $2 \times 2$  case) and output the maximum value among the input data. This temporary maximum value will be fed back to the comparator's input and regarded as one additional input to compare with the next clock cycle's input data. This procedure will be duplicated till the whole pooling window's input data is scanned. After that the output enabling signal will be validated and output the maximum value in the pooling window.

## VI. RESULTS

The accelerator was implemented in TSMC 65nm technology and the layout characteristics of the accelerator are shown in Fig. 13. The core dimension is  $2\text{mm} \times 2.5\text{mm}$  and achieves a peak throughput of 152 GOP/s at a 500MHz core clock. Since the core can support both arbitrary sized convolution layer and the pooling function, it can be used to accelerate major CNNs. A summary of the chip specifications is listed in Table II. The power is based on the synthesis report from the Synopsis Design Compile, while the area and clock speed are based on Place&Route report in Cadence. Here, PE is representing the processing engine in the chip which is a multiplier in each CU. The energy-efficiency is defined as the peak throughput divided by the dynamic power consumption.

The area breakdown of the accelerator is shown in Fig. 14. The area estimation includes the logic cells, registers,

TABLE II  
PERFORMANCE SUMMARY

Technology	TSMC 65nm RF 1P6M
Supply Voltage	1V
Clock Rate	Up to 500MHz
Dynamic Power Consumption	350mW @ 500MHz
Core Area	2mm x2.5mm
Gate Count	1.3M
Number of PEs	144
On-Chip Single Port SRAM	96 K bytes
Scratch Pads Memory	16K byte
Peak Throughput	152GOPS
Energy Efficiency	434GOPS/W
Arithmetic Precision	16-bit fixed-point
Supported CNN feature	filter kernel size:1-23
	Num. of filters:1-1024
	Num. of channels:1-1024
	Num. of features:1-1024
Supported Pooling feature	Stride Size : 1, 2, 4
	Avg Pool Size: 1-23
	Max Pool Size:2,3




Fig. 14. Area breakdown of the accelerator.

and single port/dual port SRAMs generated by the ARM compiler. As it shows, the CU engine only occupies 17% of the total area. The majority of the area is occupied by the buffer bank and scratchpad. The scratchpad is designed using dual port SRAM to support the continuous streaming while buffer bank is implemented as single port SRAM. Although the scratchpad memory size is only 1/6 compared to the buffer bank, it is still occupied more than half of the buffer bank's area.

To verify the performance of the accelerator, we have downloaded the hardware accelerator IP into the Xilinx Zynq-7200 FPGA and demonstrate the core's functions using modified LeNet-5 [19] to detect the traffic sign. The filter's weights are fetched from the DRAM through the FPGA's existing DMA controller. The DMA controller is configured as 256-depth 64-bits width. The traffic-sign net includes three convolution layers and two pooling layer and its architecture is summarized in Table III.

The application processor (AP) integrated into the FPGA is used to control the accelerator and initiate the computation. Through using the DMA controller inside the FPGA, the accelerator can successfully access the data and the weights stored in the DRAM. The demonstration setup is shown in Fig.15. The demonstration begins with downloading a traffic sign into the FPGA. After the computation, the detected traffic sign



TABLE III  
TRAFFIC SIGN CNN ARCHITECTURE

Layer	Type	Channel Size	Channel No.	Kernel Size	Stride
1	Input	32 x 32	3	—	—
2	convolution	32 x 32	3	5x5	1
3	max pooling	32 x 32	64	2x2	2
4	convolution	16 x 16	64	5x5	1
5	max pooling	16 x 16	16	2x2	2
6	convolution	8 x 8	16	5x5	1
7	fully-connected	8 x 8	16	—	—



Fig. 15. Traffic Sign Demonstration on the Xilinx Zynq-7200 FPGA.

TABLE IV  
PERFORMANCE COMPARISON

	This work	[10]	[12]
Core Area	5mm <sup>2</sup>	12mm <sup>2</sup>	16mm <sup>2</sup>
Peak Throughput	154GOPS	84GOPS	64GOPS
Gate Count	1.3M	1.2M	3.2M
Supply Voltage	1V	0.82-1.17V	1.2V
Peak Throughput	152GOPS	84GOPS	64GOPS
Energy Efficiency	434GOPS/W	166GOPS/W	1.4TOPS/W
Technology	65nm	65nm	65nm
Precision	16-bit	16-bit	16-bit
MaxPool Support	Yes	No	Yes
AvgPool Support	Yes	No	No

result will be sent back to the PC and display on the monitor. A raw video demonstration is shown in [20].

Even the demonstrated LeNet-5 Model only has an input channel size of  $32 \times 32$ , this accelerator can fit for channel size that larger than this. In fact, a large-sized channel can improve the energy-efficiency of the system. This is due to the fact that large-sized channels lead to more filter-weights reuse during the computation. For example, a  $100 \times 100$  input channel will result in approximately 10000 times filter-weights reuse during the scanning of the image, while a  $10 \times 10$  input channel only has 100 times filter-weights reuse.

When the input channel or intermediate data size is larger than the total available SRAM size. A DMA controller is needed to exchange data between DRAM and on-chip SRAM. This will cost a large energy consumption as the intermediate data is exchanged between the DRAM and the on-chip SRAM.

In addition, the data format in this hardware accelerator is set to be 16-bit fixed point to achieve minimized

hardware cost. Through re-designing the multiplier and adder in the CU block, this architecture can also be used with other data formats such as 16-bit floating point, 32-bit floating points or 8-bit fixed point.

Table IV is a comparison of the designed accelerator with other reported work. As it shows, this accelerator achieves high energy efficiency and comparable performance with low area cost, making it suitable to be integrated into the IoT devices.

## VII. CONCLUSION

In this paper, we propose a streaming architecture for the CNN hardware accelerator. The proposed accelerator optimizes the energy efficiency by reducing unnecessary data movement. It also supports arbitrary window sized convolution by using filter decomposition technique. In addition, pooling function is also supported in this accelerator through integrating separate pooling module and proper configuration of the convolution engine. The accelerator is implemented in TSMC 65nm technology with a core size of 5mm<sup>2</sup>. A traffic-sign net is implemented using this hardware IP and verified on the FPGA. The result shows that this accelerator can support most popular CNNs and achieve 434GOPS/W energy efficiency, making it suitable to be integrated with the IoT devices.

## REFERENCES


- [1] Y. LeCun, Y. Bengio, and G. Hinton, "Deep learning," *Nature*, vol. 521, pp. 436–444, May 2015.
- [2] A. Krizhevsky, I. Sutskever, and G. E. Hinton, "ImageNet classification with deep convolutional neural networks," in *Proc. Adv. Neural Inf. Process. Syst.*, vol. 25, 2012, pp. 1097–1105.
- [3] D. Silver *et al.*, "Mastering the game of Go with deep neural networks and tree search," *Nature*, vol. 529, no. 7587, pp. 484–489, Jan. 2016.
- [4] S. Bilal, *How to Classify Images With TensorFlow Using Google Cloud Machine Learning and Cloud Dataflow*. Accessed on Nov. 1, 2016. [Online]. Available: <https://cloud.google.com/blog/big-data/2016/12/how-to-classify-images-with-tensorflow-using-google-cloud-machine-learning-and-cloud-dataflow>
- [5] *Amazon Rekognition*. Accessed on Nov. 1, 2016. [Online]. Available: <https://aws.amazon.com/rekognition/>
- [6] K. Simonyan and A. Zisserman, "Very deep convolution networks for large-scale image recognition," *CoRR*, pp. 1–14, Sep. 2014.
- [7] R. Hameed *et al.*, "Understanding sources of inefficiency in general-purpose chips," in *Proc. 37th Annu. Int. Symp. Comput. Archit.*, 2010, pp. 37–47.
- [8] M. Horowitz, "Computing's energy problem (and what we can do about it)," in *IEEE Int. Solid-State Circuits Conf. (ISSCC) Dig. Tech. Papers*, Feb. 2014, pp. 10–14.
- [9] S. Han *et al.*, "EIE: Efficient inference engine on compressed deep neural network," in *Proc. 43rd Int. Symp. Comput. Archit.*, 2016, pp. 243–254.

- [10] Y.-H. Chen, T. Krishna, J. S. Emer, and V. Sze, "Eyeriss: An energy-efficient reconfigurable accelerator for deep convolutional neural networks," in *IEEE Int. Solid-State Circuits Conf. Dig. Tech. Papers (ISSCC)*, San Francisco, CA, USA, Jan./Feb. 2016, pp. 262–263.
- [11] C. Zhang, P. Li, G. Sun, Y. Guan, B. Xiao, and J. Cong, "Optimizing FPGA-based accelerator design for deep convolution neural networks," in *Proc. ACM/SIGDA Int. Symp. Field-Program. Gate Arrays*, 2015, pp. 161–170.
- [12] J. Sim, J.-S. Park, M. Kim, D. Bae, Y. Choi, and L.-S. Kim, "A 1.42TOPS/W deep convolution neural network recognition processor for intelligent IoE systems," in *IEEE Int. Solid-State Circuits Conf. (ISSCC) Dig. Tech. Papers*, Jan./Feb. 2016, pp. 264–265.
- [13] A. Shafiee *et al.*, "ISAAC: A convolutional neural network accelerator with in-situ analog arithmetic in crossbars," in *Proc. 43rd Int. Symp. Comput. Archit.*, 2016, pp. 14–26.
- [14] M. D. Pickett, "The materials science of titanium dioxide memristors," Ph.D. dissertation, Dept. Mater. Sci. Eng., Univ. California, Berkeley, Berkeley, CA, USA, 2010.
- [15] Y. Zeng, X. Xu, Y. Fang, and K. Zhao, "Traffic sign recognition using extreme learning classifier with deep convolution features," in *Proc. Int. Conf. Intell. Sci. Big Data Eng. (ISIDE)*, Suzhou, China, 2015, pp. 1–10.
- [16] A. Krizhevsky, I. Sutskever, and G. Hinton, "ImageNet classification with deep convolution neural networks," in *Proc. Conf. Workshop Neural Inf. Process. Syst.*, 2012, pp. 1106–1114.
- [17] K. He, X. Zhang, S. Ren, and J. Sun, "Deep residual learning for image recognition," in *Proc. IEEE Conf. Comput. Vis. Pattern Recognit. (CVPR)*, Las Vegas, NV, USA, Jun. 2016, pp. 770–778.
- [18] S. Gupta, A. Agrawal, K. Gopalakrishnan, and P. Narayanan, "Deep learning with limited numerical precision," in *Proc. 32nd Int. Conf. Int. Conf. Mach. Learn. (ICML)*, vol. 37. Lille, France, 2015, pp. 1737–1746.
- [19] Y. LeCun, L. Bottou, Y. Bengio, and P. Haffner, "Gradient-based learning applied to document recognition," *Proc. IEEE*, vol. 86, no. 11, pp. 2278–2324, Nov. 1998.
- [20] Kneron Inc. *Hardware IP Demo*. Accessed on Jun. 1, 2017. [Online]. Available: <https://www.youtube.com/watch?v=ttdSLXmBEWE>



the Henry Samueli Fellowship in 2012 and the Broadcom Fellowship in 2015.

**Yilei Li** received the B.S. and M.S. degrees in microelectronics from Fudan University, Shanghai, China, in 2009 and 2012, respectively, and the Ph.D. degree in electrical engineering from the University of California at Los Angeles, Los Angeles, CA, USA, in 2016. He is currently with Novumind Inc., where he is involved in the hardware development. His current research interests include circuit and system design for emerging applications, including software-defined radio, multiband RF interconnect, and AI acceleration hardware. He was a recipient of




**Junjie Su** received the B.S. degree in electrical engineering and computer science from the University of California at Berkeley in 2006, and the M.S. degree in electrical and computer engineering from the University of California at San Diego in 2008. He has over seven years ASIC design and verification industry experience and has worked in many world-class companies, such as Marvell, Broadcom, and Synopsys. He is currently a Principal A.I. Research Scientist with Kneron Inc.



**Li Du** (M'16) received the B.S. degree in information science and engineering from Southeast University, Nanjing, China, in 2011, the M.S. degree majoring in electrical engineering from the University of California at Los Angeles (UCLA), and the Ph.D. degree from UCLA in 2016. At UCLA, he was with the High-Speed Electronics Laboratory and in charge of designing high-performance mixed-signal circuits for communication and touch-screen systems. In 2012, he was an Intern with the Broadcom Corporation FM Radio Team and in charge of

designing the second-order continuous-time delta-sigma ADC for directly sampling FM radios. From 2013 to 2016, he was with Qualcomm Inc., designing mixed signal circuits for cellular communications. He is currently a Hardware Architect with Kneron Inc.



**Yuan Du** (M'17) received the B.S. degree (Hons.) in electrical engineering from Southeast University, Nanjing, China, in 2009, and the M.S. and Ph.D. degrees in electrical engineering from the University of California at Los Angeles in 2012 and 2016, respectively. He has co-authored several leading journal and conference papers, such as JSSC, MTT, TVLSI, TCAS I&II, TCAD, ISSCC, VLSI, CICC, and IMS. His research interests include designs of domain-specific computing hardware accelerator, high-speed wireline/SerDes, and RFICs. He was a

recipient of the Microsoft Research Asia Young Fellowship in 2008, the Southeast University Chancellor's Award in 2009, and the Broadcom Fellowship in 2015.



**Yen-Cheng Kuan** (M'12) received the B.S. degree from National Taiwan University, Taipei, Taiwan, and the M.S. and Ph.D. degrees from the University of California at Los Angeles, Los Angeles, CA, USA, all in electrical engineering. From 2004 to 2007, he was a System Engineer with Realtek Semiconductor Corporation, Irvine, CA, USA, where he was involved in the design of ultrawideband system-on-a-chip. From 2009 to 2016, he was a Research Scientist with HRL Laboratories, Malibu, CA, USA, where he was involved in the designs of software-

defined radios, compressed-sensing receivers, multirate signal processors for high-speed ADCs, and image processors. He is currently the MediaTek Junior Chair Professor with the International College of Semiconductor Technology, National Chiao Tung University, Hsinchu, Taiwan. He holds over 20 U.S./international granted/applied patents. His current research interests include low-power architecture design and system level design for various DSP applications and communication systems. He was a recipient of the HRL New Inventor Award for his invention contribution.



**Chun-Chen Liu** received his B.S. degree in Electrical Engineering from the National Cheng Kung University, Taiwan, in 2003, and the M.S. degree from UCSD with a co-pai of UC Berkeley in 2008 and his Ph.D degree from UCLA in 2017. From 2007 to 2010, He was the Technical Officer with Wireless Info Tech Ltd. (acquired by VanceInfo, NYSE:VIT). He has served as the technical team leader, and managing and leading several research and production development in Samsung, Mstar, and Qualcomm. With his unique entrepreneurial and

managerial expertise, he has successfully founded and co-founded several startups, including Skyvin and Rapidbridge (acquired by Qualcomm in 2012). He is currently the CTO of Kneron, San Diego, CA, USA. He was a recipient of the IBM Problem Solving Award based on the use of the EIP tool suite in 2007. Two of his papers were nominated as the best paper award candidates at the IEEE/ACM International Conference on Computer-Aided Design in 2007 and the IEEE International Conference on Compute Design in 2008.



**Mau-Chung Frank Chang** (M'79–SM'94–F'96) received the B.S. degree from National Taiwan University in 1972, the M.S. degree from National Tsing Hua University in 1974, and the Ph.D. from National Chiao Tung University in 1979. He served as the Chair of the EE Department from 2010 to 2015. He is currently the President of National Chiao Tung University, Hsinchu, Taiwan, and also the Wintek Distinguished Professor of electrical engineering with the University of California at Los Angeles (UCLA), Los Angeles, CA, USA.

Before joining UCLA, he was the Assistant Director of the High Speed Electronics Laboratory, Rockwell International Science Center, from 1983 to 1997, Thousand Oaks, CA, USA. In this tenure, he developed and transferred the heterojunction bipolar transistor (HBT) integrated circuit technologies from the research laboratory to the production line (later became Skyworks). The HBT productions have grown into multi-billion dollar businesses and dominated the cell phone power amplifier and front-end module markets for the past 20 years (currently exceeding ten billion-units/year and 50 billion units in the last decade). Throughout his career, he has pursued his research in areas of high-speed electronics, integrated circuit/system designs for radio, radar and imagers, and multiband interconnects for intra- and inter-chip communications. He is recognized by his memberships with the U.S. National Academy of Engineering in 2008 and the Academia Sinica of Taiwan in 2012. He received the David Sarnoff Award in 2006.